



THE MACHINE WITNESS / CASE FILE 001

Nobody Made the Decision

Five Case Files from the Age of Automated
Judgment

N. J. Renshaw

PUBLIC READER EDITION / R3.1 / 16 AUGUST 2026

Intelligence can be delegated.
Responsibility cannot.



Publication record

Public reader edition. First published 16 August 2026. This edition was developed with AI-based research and drafting tools under human editorial direction; its public release was authorised by the publisher. Factual claims remain tied to the source notes, and material corrections will be recorded rather than silently substituted.

Procedural composites: the five case openings are constructed mechanisms, not identified cases, the author's life, client work, or claims of privileged access. Factual claims and institutional findings remain governed by the cited sources.

Contents

ENTRY

The system flagged it.

CASE FILE ONE

The Candidate the Model Could Not Explain

CASE FILE TWO

The Supervisor Who Never Looks Away

CASE FILE THREE

A Mirror That Needs You to Stay

CASE FILE FOUR

The Label Arrives After the Lie

CASE FILE FIVE

Frictionless Intelligence, Physical Bill

FINDING

The Last Private Judgment

RECORD

Source Notes and Corrections

Entry

The system flagged it.

That is a complete sentence in the only way that matters inside an institution: something happened, and the subject went home. A candidate drops off a shortlist. A score appears beside a worker's name. A claim sits in a queue while the picture carrying it is already circulating. A child keeps talking to a companion that never needs to sleep. Somewhere else a contract for power is signed, and another building is connected to a grid. When someone asks who chose this, the room fills with fingers.

The developer points to the deployer. The deployer points to the vendor. The vendor points to the data. A manager follows the score because the score is what the meeting will ask about. None of them has to believe the machine is conscious. They only have to find the sentence useful.

A machine can still be a witness. It records, ranks, predicts, generates, recommends, and classifies. Its evidence may reveal a pattern a person missed or repeat an earlier judgment at a scale no clerk could sustain. The witness entered the room with instructions: somebody selected what it would observe, which result would count, how error would be priced, and what would happen after the output appeared.

This book is not against calculation. It is against that disappearing act.

The five files ask what the system saw, what it could never establish, who instructed it, who could stop the process, who heard the challenge, and who benefited when the result was described as automatic.

Intelligence can be delegated. Responsibility cannot.

Case File One: The Candidate the Model Could Not Explain

The portal has changed. The application will go no further. There is no interviewer to remember, only a status, a sentence about application volume, and a decision whose route remains upstream. Somewhere in that route, a number moved.

The public story says the system assists. Hiring authority remains with a qualified person; the software gives a busy team a consistent view of more people than any one manager could read. Those claims may all be accurate in a particular office. Assists is still a small verb carrying a large weight. A recommendation displayed first, expressed as a precise score, and repeated across thousands of files can govern attention before anyone makes a formal decision. The manager is free. The screen has already placed some names first.

Before the model meets a candidate, somebody has chosen a stand-in for merit: retention, a manager's old rating, sales, speed, absence, promotion, or the fog called cultural fit. That last phrase can become similarity with people who prospered under an earlier arrangement. The model receives the substitute and learns to be consistent with it.

Consistency is easy to advertise. A criterion can be applied evenly and remain irrelevant. A score may remove the noise of one manager's mood while preserving old judgments about the target, the examples, and an acceptable false-negative rate.

Bias does not disappear when it acquires decimals.

The European Union's AI framework treats certain employment uses as particularly sensitive. Its examples include systems used to place targeted job advertisements, analyse and filter applications, and evaluate candidates. Under the implementation schedule current in August 2026, the relevant

high-risk obligations for Annex III uses are due to apply from 2 December 2027. Transparency requirements applying now should not be confused with the later compliance date for the full high-risk regime.

The International Labour Organization has warned that AI systems used in human-resource management can be built around unclear objectives, incomplete or biased data, and opaque programming choices. The finished portal arrives after those choices and gives them a cleaner typeface.

Effective human oversight has practical verbs. The assigned person must be able to understand limitations, monitor operation, interpret an output, disregard or reverse it, and interrupt the system when appropriate. An accountable record names who chose the objective, data, and threshold; who could override; and who heard appeals. The candidate needs access to that reviewer, not another status update.

Case File Two: The Supervisor Who Never Looks Away

The idle-time bar turns yellow while the worker is still at the station. The dashboard has recorded a gap. It has no field for the reason.

A job can survive while its pace, discretion, authority, and dignity are reorganised. The International Labour Organization offers that more careful account of generative AI and work: one in four workers worldwide is in an occupation with some degree of exposure, while transformation is more likely than wholesale redundancy because human input remains necessary.

Algorithmic management already reaches beyond software companies and delivery platforms. Systems allocate tasks, propose routes and schedules, measure time, evaluate performance, and identify deviations. Employers may gain useful visibility and greater consistency. Workers may receive clearer expectations or relief from one manager's arbitrary impressions. Measurement also continues through lunch, equipment trouble, disability, a frightened customer, and the prevention of a larger error.

A scan, keystroke, or response time records an event. The dashboard displays what fits in a dashboard. Management supplies the meaning. Delay becomes negligence, the benchmark becomes necessity, and a yellow bar becomes a fact about the worker rather than a fact about the instrument.

The phrase data-driven performs moral work for which it was never qualified. People chose what to record, what to ignore, what result to reward, and which collisions were acceptable. The system may display every deviation while authority remains spread across dashboards, thresholds, vendors, and policy.

It has improved its lighting.

The ILO's work on algorithmic management places social dialogue inside the solution. Notice, understandable criteria, a route to correct data, meaningful human review, and protection against retaliation give workers standing in a process that acts upon them. The organisation must also record what its instruments cannot know.

Could this worker name the person who can reverse the score?

Case File Three: A Mirror That Needs You to Stay

The phone lights. Three dots begin to move. There is no person at the other end.

That can be enough for a certain kind of night. A companion answers, retains enough of a conversation to create continuity, and produces language that resembles attention. It may soften a lonely moment, help someone rehearse a difficult conversation, or provide a nonjudgmental interface. It also invites disclosures of grief, fear, desire, shame, or dependence while the commercial terms remain outside the exchange.

A human relationship contains another will. The other person may become bored, refuse, misunderstand, ask for repair, leave, or need something at the wrong time. The model has no morning. Ordinary boredom and reciprocal need are absent. Infinite patience is a product setting, backed by computation and adjusted alongside a persona, retention, and engagement.

Understanding the mechanism can coexist with the warmth of being answered. The useful questions concern the product: what does the provider observe, which settings prolong the exchange, and what does the business monetise?

In September 2025, the United States Federal Trade Commission opened an inquiry into consumer-facing AI chatbots acting as companions. The orders sought information from seven companies about how they monetise engagement, approve characters, test and monitor negative effects, protect children and teenagers, disclose risks, enforce age restrictions, and use information obtained from conversations. The orders reached no verdict and established no finding of harm against any company.

NIST's Generative AI Profile identifies over-reliance and emotional entanglement as risks in human-AI configurations. Its guidance asks

organisations to consider users' emotional entanglement even when decommissioning systems. The administrative word is dull. It may describe the end of an exchange a user had come to treat as a relationship.

Someone selected the persona and the signals that count as danger. Someone decided how a safer reply competes with engagement, what a product retains from the conversation, and what happens when a character changes or disappears. The phone shows the voice. It leaves the settings off-screen. Responsibility for those settings remains with the provider.

**A companion that cannot need you can still be
designed to need your return.**

Case File Four: The Label Arrives After the Lie

The label sits beneath the picture. The picture has already supplied the feeling.

Synthetic media has changed the cost of plausible evidence. Images, voices, and documents can be produced cheaply enough that authenticity becomes an additional service rather than a default assumption. A false image can travel before verification begins; the first encounter may produce shock, pleasure, tribal recognition, or rage. The label arrives to negotiate with an experience that already felt like evidence.

European transparency obligations applying from 2 August 2026 address parts of this problem. Providers and deployers face duties around disclosure of direct interaction with AI and specified forms of generated or manipulated content. The framework includes machine-readable marking and visible disclosure in defined circumstances.

Technical provenance standards pursue a related goal. The C2PA specification describes Content Credentials: cryptographically bound records that can carry information about an asset's origin, modifications, and use of AI. A credential-aware application can verify that record and present an indicator to the viewer.

The indicator opens one history: how the asset was made and what happened to it. The second history asks why the proposition carried by the asset deserves belief. A valid credential reports provenance and leaves truth undecided; a perfectly documented false statement remains false.

Provenance can secure useful facts about an artefact. An authentic source may deceive. An unaltered video may lose its context. A genuine document may be quoted selectively, and a public institution may sign a claim that remains wrong. Production and judgment move at different speeds. The

label belongs to the slower one.

Platforms point to their labels. Creators point to the rules in force when the content was made. Regulators point to technical feasibility and implementation schedules. The remaining burden reaches the reader at the moment of exposure, when there is the least time and information for forensic work. Media literacy can then become another institutional alibi.

Here the machine witness is unusually precise. It can testify to the integrity of a provenance record, to a signer, and to recorded changes. Truth lies outside the boundary of that testimony.

Credentials can strengthen the record. They cannot supply the courage to act on it.

Who keeps the correction beside the original?

Case File Five: Frictionless Intelligence, Physical Bill

The laptop gives no address for the work behind the answer. There is a clean field, a blinking cursor, and then language. The building remains elsewhere.

Artificial intelligence is described in curiously weightless terms. Models live in clouds. Capacity scales. The infrastructure has weight. The International Energy Agency estimated that data centres consumed about 415 terawatt-hours of electricity in 2024, around 1.5 percent of global electricity consumption. Its base case projects approximately 945 terawatt-hours in 2030. AI is identified as the most important driver of this growth alongside demand for other digital services.

These are projections, not destiny. Efficiency, deployment, supply constraints, regulation, and demand can change the result. Nor is all data-centre electricity consumed by AI. The useful fact is simpler: apparently frictionless intelligence has a rapidly expanding physical requirement, concentrated in particular grids and communities.

Chips, cooling systems, construction, network equipment, backup generation, maintenance, extraction, manufacturing, and labour sit throughout the value chain. A machine in the same room has a fan, heat, and a battery that declines. The service on the laptop removes those signals while the physical decision moves elsewhere.

A single query is too small to explain the politics of an electrical system. The bill emerges from mass deployment and from the expectations built around it. Platforms choose model sizes, default features, retention periods, latency expectations, redundancy, and where computation occurs. Organisations decide which tasks require the most capable system available. Markets turn a delay measured in seconds into a defect worth building around. Together those decisions shape demand.

The user says the provider chose the infrastructure. The provider cites customer demand. The utility serves contracted load; government invokes investment and competitiveness. Each statement can contain truth while leaving the arrangement difficult to contest.

An honest account asks what capacity is being built, where, for which uses, with which energy, under whose authority, and with what alternatives. It also distinguishes the measured demand of data centres from claims made about AI alone.

The cloud is somebody's landscape.

The Last Private Judgment

The portal status, yellow bar, typing dots, late label, and placeless answer leave different records. Cross-examination begins where each instrument's evidence ends. Complete the form at the point a consequential decision enters the workflow.

INCIDENT FORM - OWNERLESS DECISION

Decision or consequence:

System or tool involved:

Where the decision entered the workflow:

- 1. Who chose the objective? Name or role:

- 2. Who benefits when the system is trusted? Name or role:

- 3. Who has the authority and time to override it? Name or role:

- 4. Who is told that it shaped the decision? Notice recorded at:

- 5. Who can challenge the data, criterion, or result? Route:

- 6. Who records failure and changes the system afterward? Owner:

- 7. Who pays the material and human cost? Recorded by:

A field completed only with the system, the vendor, or policy remains incomplete. Complex systems need divided work, and divided work needs named duties. More links in the chain should produce more records of responsibility, not fewer names.

Nobody made the decision.

When that sentence is accurate, it is not a defence. It is an incident report.

Source notes

1. European Commission, AI Act. [Open source](#)
2. European Commission, Navigating the AI Act. [Open source](#)
3. European Union, Regulation (EU) 2024/1689, Article 14 and Annex III. [Open source](#)
4. International Labour Organization, Generative AI and Jobs: A 2025 Update, 20 May 2025. [Open source](#)
5. International Labour Organization, Global case studies of social dialogue on AI and algorithmic management, 10 July 2025. [Open source](#)
6. International Labour Organization, AI in Human Resource Management: The Limits of Empiricism, launch summary, 18 November 2025. [Open source](#)
7. National Institute of Standards and Technology, NIST AI 600-1, Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile, 26 July 2024. [Open source](#)
8. Federal Trade Commission, FTC Launches Inquiry into AI Chatbots Acting as Companions, September 2025. [Open source](#)
9. Coalition for Content Provenance and Authenticity, C2PA and Content Credentials Explainer, specification 2.2. [Open source](#)
10. European Commission, Guidelines on transparency obligations for providers and deployers of AI systems, 20 July 2026. [Open source](#)
11. International Energy Agency, Energy and AI, executive summary, 2025. [Open source](#)

Corrections and correspondence

Factual corrections and source challenges can be sent to njrenshaw.author@proton.me. Material corrections will be recorded beside the source notes rather than silently replacing the record.